

Tutorial 1: Linking and Loading

Model Solutions

Spring 2023

Question 1

Use generate script to generate elf files for the given programs.

```
Disassembly of section .text:

0000000000000000 <main>:
#include<stdio.h>
void swap();
int buf[2] = {34,56};
int main(){
    0:  f3 0f 1e fa          endbr64
    4:  55                   push    %rbp
    5:  48 89 e5             mov     %rsp,%rbp
    swap();
    8:  b8 00 00 00 00      mov     $0x0,%eax
    d:  e8 00 00 00 00      callq   12 <main+0x12>
    printf("buf[0]=%d buf[1]=%d\n",buf[0],buf[1]);
    12:  8b 15 00 00 00 00    mov     0x0(%rip),%edx
    18:  8b 05 00 00 00 00    mov     0x0(%rip),%eax
    1e:  89 c6               mov     %eax,%esi
    20:  48 8d 3d 00 00 00 00 lea     0x0(%rip),%rdi
    27:  b8 00 00 00 00      mov     $0x0,%eax
    2c:  e8 00 00 00 00      callq   31 <main+0x31>
    return 0;
    31:  b8 00 00 00 00      mov     $0x0,%eax
}
```

Relocation section '.rel.text' at offset 0xb68 contains 5 entries:

Offset	Info	Type
Symbol's Value	Symbol's Name + Addend	
000000000000000e	0000001200000004 R_X86_64_PLT32	
0000000000000000	swap - 4	
0000000000000014	0000000f00000002 R_X86_64_PC32	
0000000000000000	buf + 0	
000000000000001a	0000000f00000002 R_X86_64_PC32	
0000000000000000	buf - 4	
0000000000000023	0000000500000002 R_X86_64_PC32	
0000000000000000	.rodata - 4	
000000000000002d	0000001300000004 R_X86_64_PLT32	
0000000000000000	printf - 4	

- A portion of the file main.elf file is shown in the bottom figure. It shows the locations needing relocation in the text segment. Specifically it shows that locations e, 14, 1a, 23 and 2d need

relocation. No location in the data segment needs relocation.

- It also says that the relocation type is `R_X86_64_PLT32/PC32`. These are PC-relative relocations, and the value to be relocated is a 32 bit quantity in each case.
- Turning to the `main.obj` file on the top: The value to be relocated at `e` stands for the procedure `swap`.

If *loc* is a location before relocation, let *final(loc)* denote the final address of *loc* after relocation. Let *offset* be the value that is filled in the 4 bytes starting from `e`. Now the fact that relates these values is that when the call at `d` is executing, the PC is pointing to the next instruction *final(12)*, and the address of `swap` is obtained by adding *offset* to PC.

$$\begin{aligned} \text{PC} + \text{offset} &= \text{loc}(\text{swap}) \\ \text{offset} &= \text{loc}(\text{swap}) - \text{PC} \\ \text{offset} &= \text{loc}(\text{swap}) - (\text{loc}(\text{e}) + 4) \\ \text{offset} &= \text{loc}(\text{swap}) - \text{loc}(\text{e}) - 4 \end{aligned}$$

- After the linker has determined the values of *loc(swap)* and *loc(e)*, it uses the above formula to obtain the value of *offset*. The value -4 is what is called **Addend** in `main.elf`.

`swap.c` and `others.c` can be analyzed in a similar manner.

There is one more point that needs attention: In `swap.c` there is a location in the data area that needs relocation. This can be seen as an entry in the data relocation area in `swap.elf`

```
Relocation section '.rela.data.rel' at offset 0x710 contains 1 entry:
  Offset          Info          Type           Symbol's Value  Symbol's Name + Addend
0000000000000000  000000011000000001 R_X86_64_64     0000000000000000 buf + 0
```

This is an absolute relocation. The final address of the target (and not an *offset*) is put directly at the place of relocation. And since this is an address, it is a 64-bit quantity.

Question 2

Annotate fragments of target code with the source statements that they correspond to. Annotate each local variable and parameter with its (relative address).

test1.c

<pre> 1 int main() 2 { 3 int a=1, b=1; 4 while (a<=10) 5 { 6 b=b*a; 7 a++; 8 } 9 return b; 10 }</pre>	<pre> 1 main: 2 pushq %rbp 3 movq %rsp, %rbp 4 movl \$1, -8(%rbp) — b=1 5 movl \$1, -4(%rbp) — a=1 6 jmp .L2 — while condition 7 .L3: 8 movl -4(%rbp), %eax — Copy a 9 imull -8(%rbp), %eax — a*b 10 movl %eax, -4(%rbp) — b=b*a 11 addl \$1, -8(%rbp) — a++ 12 .L2: 13 cmpl \$10, -8(%rbp) — a<=10 14 jle .L3 — to inside while loop 15 movl -4(%rbp), %eax — value of b 16 popq %rbp — return b 17 ret</pre>
--	--

Only two local variables, *a* is in $-4(\%rbp)$, or offset -4 , *b* is in offset -8 .

test2.c

<pre> 1 struct data{ 2 int sum; 3 int b[5]; 4 }; 5 6 int main() 7 { 8 struct data rec1; 9 rec1.sum=0; 10 rec1.b[0]=2; 11 rec1.sum=rec1.sum+ 12 rec1.b[0]; 13 return rec1.sum; 14 }</pre>	<pre> 1 main: 2 pushq %rbp 3 movq %rsp, %rbp 4 movl \$0, -32(%rbp) — rec1.sum=0 5 movl \$2, -28(%rbp) — rec1.b[0]=2 6 movl -32(%rbp), %edx — rec1.sum 7 movl -28(%rbp), %eax — rec1.b[0] 8 addl %edx, %eax — sum + b[0] 9 movl %eax, -32(%rbp) — rec1.sum 10 movl -32(%rbp), %eax — value of rec1.sum 11 popq %rbp — return rec1.sum 12 ret</pre>
--	--

The local variable `rec1` contains `sum`, which is stored in offset `-32`, the array `b` spans from offset `-28` to `-12`.

Question 3

- a) The `main` in `module1`. The `main` in `module2` is an uninitialized global and therefore a weak symbol.
- b) `gcc` will give an error because both `main`s are strong symbols.
- c) No clash of symbols. `main` in `module2` is not visible outside of `module2`.

Question 4

User stack: `p`

Heap: `*p`

Read only: the functions `main` and `f`, the format string `"%d\n"`.

Read write: `x`, `y`, `a`, `k`

Question 5

The answer is:

`{14, 3}`


`{3, 14}`

Note that the `recs` in both modules default to the `rec` in `module1`. Now for `module2`, `x` is the second field and `y` is the first field of whatever `rec` resolves to. Thus `{14,3}`. The printing of `{3,14}` is obvious.

Question 6

The line `int z = a[5]` is a compile time error. `a[5]` is not a compile time constant. However, `int *x = &(a[3])` is ok since `&(a[3])` is a compile time constant.

Now just before the program starts executing, here is a (partial) view of the global memory from the point of view of the two modules:



After the execution of `c = 100.0`, this is what happens:



And since the call to `fn()` returns the `d` part of this memory, it is highly improbable that it is equal to 5.