

# COMPUTER ARCHITECTURE

## \* PERFORMANCE METRICS

\* Performance  $\propto \frac{1}{\text{Time}}$

$$\text{Time} = \underbrace{\# \text{ instructions}}_{\text{compiler tech.}} \times \underbrace{\text{cycles per instruction}}_{\text{CPI, PLD}} \times \underbrace{\text{time per cycle}}_{\text{clock speed}}$$

\* ISA (Instruction set architecture) is the interface between software and hardware. It provides a behavioural abstraction to the software and a set of specifications to the hardware.  
Eg:- ARM ISA can run on Apple M1 or Snapdragon <sup>or ARM native</sup> hardware.  
x86 ISA can run on Intel or AMD hardware.

\* The hardware i.e. computer organisations is chosen based on its frequent job, i.e. floating point ops, integer ops, data i/o.

We generally rate machines based on:-

- No. of <sup>Cores</sup> ~~cores~~ (For multithreaded programs)
- GPU (For vectorized operations)
- Clock speed
- RAM

\* SPEC (System Performance Evaluation Consortium):

Created a set of programs representative of all programs.

These are the SPEC CPU Benchmarks (2017).

Companies generally release the ratio of time taken by their machine to the time taken by a standard machine.

\* Another metric used to evaluate is MIPS (Million instructions per second).  $= \frac{\# \text{ Instr.}}{\text{Time} \times 10^6}$

$$= \frac{Ic}{Ic \times CPI \times \frac{1}{f} \times 10^6} = \frac{f}{CPI \times 10^6}$$

Generally used to report the upper bound of a machine's performance.



Similarly, we also have MFLOPS (Million floating point ops per sec.), generally used to measure performance of super computers.

\* To find average values, we use :-

- for time : Arithmetic mean
- for time ratios : Geometric mean
- for MIPS / MFLOPS : Harmonic mean.

### \* ISA DESIGN :

The ISA needs to perform a minimal set of arithmetics & logical operations taking convenience into account. It should also be able to ~~des~~ communicate with memory (RAM). It should also have instructions for changing control flow.

These provide enough generality to be Turing complete.

A typical instruction will look like :

Operation    Operands    Destination.

The size of the instruction will be (in bits).

$$= \lceil \log_2 (\text{No. of possible operations}) \rceil + (\text{No. of operands} + \text{dest}) \times \lceil \log_2 (\text{size of memory in bytes}) \rceil$$

Note: Running a program with limited RAM: If the size of the program is larger than the RAM size, it is run sequentially, taking only parts out of the secondary storage to the RAM. This has a lot of time overhead.

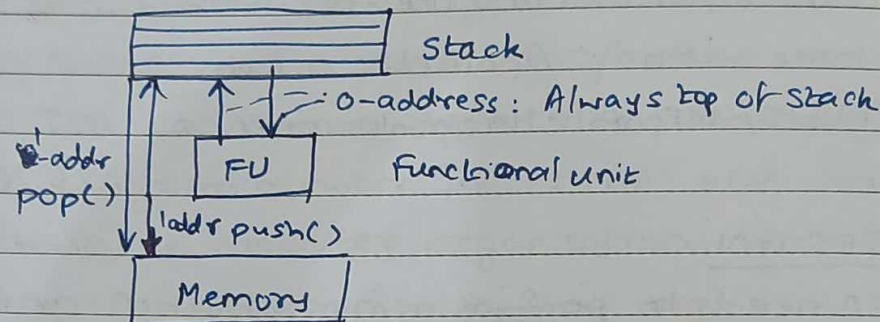
Since memory is expensive, we try to make the ISA as memory efficient as possible. Another constraint to the ISA is slowness of memory (multiple clock cycles needed).

The ISA should also address ease of programming.

### \* Targetting limited memory

In the instruction, we can reduce no. of bits needed for the addresses of the operands.

→ Can we construct a machine where all addresses are implicit i.e. only operator needs to be specified? Yes!



Eg: to calculate  $A = B + C * D$  ⇒ Convert to postfix

push C → 1 addr

$A = BCD * +$

push D → 1 addr

mul → 0 addr

push B → 1 addr

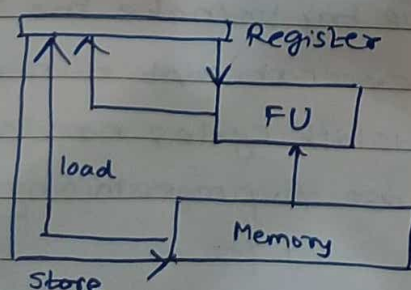
add → 0 addr

pop A → 1 addr.

If all operators are binary, then, in the parse forest: the 3 addr format takes  $3 \times (\text{no. of internal nodes})$  memory accesses. The 0 addr format takes  $(\text{No. of leaves} + \text{No. of roots})$  memory accesses.

In this stack based architecture, stack is the bottleneck.

→ 1-address format:



$A = B + C$

Load B

Add C

Store A



$A = B + C * D$  : load C ; mul D ; ~~add~~ add B ; store A ;

This is called an accumulator based architecture.

$A = B * C + D * E$  : load B ; mul C ; store ~~T~~ T<sub>1</sub> ;  
 load D ; mul E ; add T<sub>1</sub> ; store A ;

This illustrates that the accumulator is again a bottleneck.

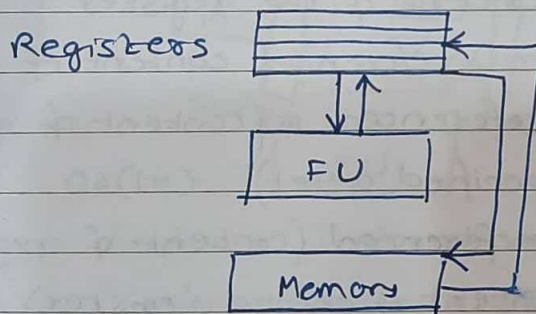
→ We can overcome the issue using a small number of registers to act as temporary storage (which also overcomes slowness of memory).

Op | AM | Addr.

↑  
 Addressing Mode: Register or Main memory.

With the introduction of registers, the accumulator was eliminated. Now, ~~add~~ arithmetic & logical operations were performed only on registers. Since ~~see~~ no. of registers

is small, we ~~could~~ <sup>can</sup> afford to have 3 address format for arithmetic and logical operations.



Thus, a 3 address format was introduced on registers.

Op | Reg1 | Reg2 | Reg3

~~to~~ To transfer data between reg. and memory,

Op | Reg | Mem.addr

$A = B + C$  : Load r1, B ; load r2, C ; add r1, r2, r0 ; store r0, A ;

If the number of registers is too small, we may need too many memory references. ~~We~~ We can ~~overwrite~~ <sup>overwrite</sup> registers after that variable has died. Based on empirical evidence, we use 32 or 64 registers. If the number of live variables is more, we may have to load & store from memory.



\* We have dealt with accessing operands from memory. But, we still need to access instructions from memory. This now becomes the bottleneck. This ~~more~~ demands more complex instructions, so that less no. of instruction reads are needed.

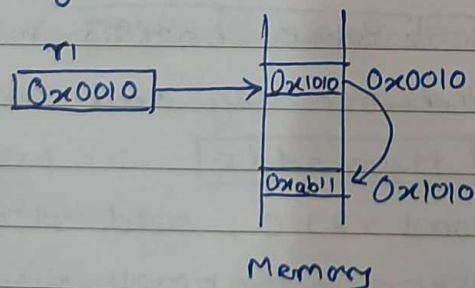
### \* Ease Of Programming:

- Pointers: We can store addresses in registers just as we store data. This also needs a dereference operator.

~~Therefore~~ This requires an addressing modes, to specify if we are referring to the content of the register or if it is the data at the address which is content of the register.

Addressing modes:

1. Direct Memory: Direct constant address:  $0x0ab1$
2. Register Direct: Content of the register:  $r1$
3. Register Indirect: Dereferenced content of register:  $(r1)$
4. Base + offset: Dereferenced (content of register + specified offset):  $(r1)40$
5. Base + index: Dereferenced (content of register + content of second register):  $(r1)(r2)$
6. Memory indirect:



7. Auto increment/decrement: Used to implement stack: Eg: `push %eax %rsp`.


Means put content of `eax` register into address contained in `rsp`, and increment the content of `rsp` suitably to point to next available language.



\* 16 → boundaries

\* Youtube: CPU on a breadboard

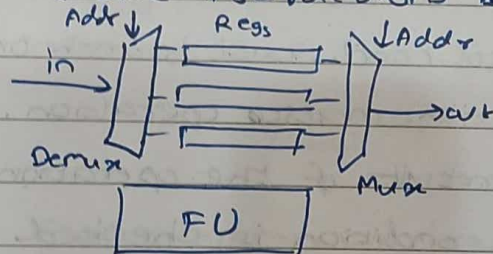
8. Constants: One of the operands is a constant. We can give the value of the constant in the instruction itself. Called Immediate Addressing. There is a limit (usually 32 bits) to the size of constant allowed; beyond that, it has to be treated as a variable.



\* Memory is byte-addressable, so the instructions are made in multiples of 8 bits.

### \* RISC ARCHITECTURE:

The machine can only use register contents for ALU, and directly use the instructions as machine signals

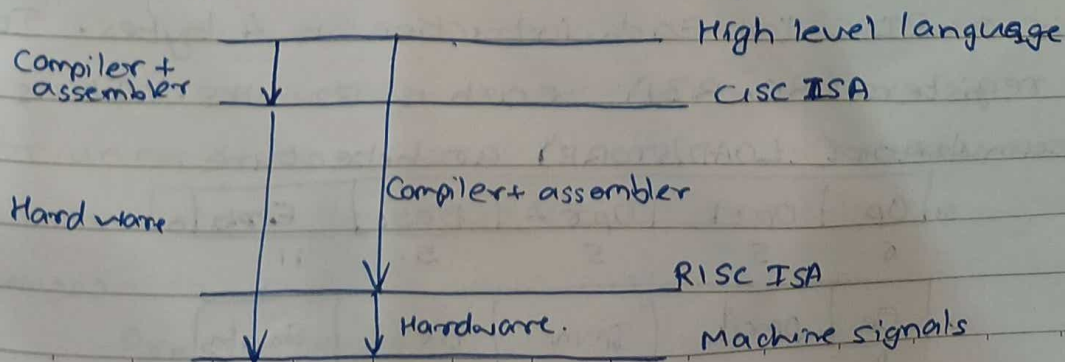


RISC = Reduced Instruction Set Computer. Eg:- ARM

CISC = Complex Instruction Set Computer. Eg:- x86

(used in RISC)


\* Type of encoding: If fixed length<sup>1 bit</sup> encoding is used, it leads to ~~lot~~ of inefficiency in CISC. Thus, variable length encoding is used by CISC. This poses problems with decoding and parallel processing of instructions.





(Note: During development we use low compiler optimization as we want it to compile ~~and~~ quickly for us to debug. ~~the~~ In production, we use high level of optimization).

Intel provides a CISC interface to its software. But internally, it converts it to a RISC pseudo-ISA. This is called ISA indirection.

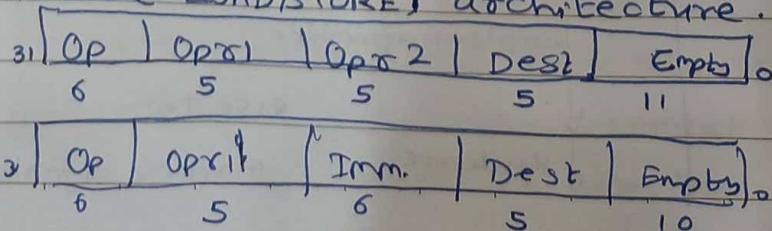


#### \* Parts of ISA:

- Arithmetic & logical
- Memory communication
- Control flow change: Unconditional branch like in functional calls. Also, conditional branch on the result of some arithmetic operation. It can happen that the result of the operation gets erased before the condition is checked. Hence, we have CPU flags, which are set when some condition is met. Eg:- carry is 1, result is 0, overflow flags. These flags are updated after every ALU operation. When we perform a jump, the final address may be absolute or relative to current address.


#### \* MIPS ISA:

It is a RISC ISA. Each instruction is 4 bytes. It has 32 registers (R0-R31), each is 32 bits. It is an ~~accumulator~~ (LOAD/STORE) architecture.





- All ALU ops. are performed only on registers.  
 It also allows immediate (constants) addressing mode. To distinguish these, we have different instructions for the immediate variants of the operation. Alternatively, the last 6 bits of the instruction are also used; for register operation, first 6 bits are 0 and last 6 are op-code; for immediate operation, first 6 bits are used for the op-code.
- Memory communication: It has only LOAD and STORE instruction. The address can be given as Base+Displacement or Base+Index.




This does  $R2 = *(R1 + 80)$  if Op = Load.

and  $*(R1 + 80) = R2$  if Op = Store.

Note: Everything here is byte addressed.


We have instructions: SB, LB, SH, LH, SW, LW  
                                     Store byte      Store Half      Store Word  
                                     /Load            /Load word    /Load

For Base+Index, we have:



- Jump: Conditional

BEGZ Reg, Loc (or BNE Reg, Loc)




as each instruction is 1 word long.

If Reg = 0, then shift Imm ~~bytes~~ words ahead.

(Note that for this, we need a program counter register, which is actually a pointer to the instruction address in RAM).

For unconditional: J Loc



Always,

Program Counter (Bytes) += Imm \* 4



We also have JAL instruction, which loads the program counter to R31 before jumping. This allows function calls. To return, we use JR Reg.

- \* When we do a jump, the Imm field gives no. of instructions to go relative to the current instruction. But in BEQ / BEQZ, it is w.r.t. next instruction.

Eg:-

SEB R3, R3	←	Set R3 to 0 on equality of R1 and R2
BEQ R3, 12	←	If R3 = 0, Go to (12+1) × 4 instructions ahead.

A jump instruction takes a 26 bit Imm field. It can move  $2^{26} \times 4$  bytes away. For a bigger jump, Multiple jumps must be taken.

Also, we have JR Reg. instruction, which says "jump to absolute location" given in register Reg. (in words).

- \* Function calls :- JAL Loc : Jump to Loc (Relative) and store (current PC + 4) in R31 to allow for return. We also have JALR Reg.

To return, we can then use JR R31.

Note that a recursive function call would over-write

R31. Hence, a stack needs to be implemented in memory.

To load a constant value into a register, we use

LHI Reg, Imm<sup>16</sup>. The Imm is moved to the most significant 16 bits of Reg. Thus to store a 32 bit address, we can load the <sup>most</sup> significant 16 bits first, then add the least significant 16 bits. We also have read-only register R0 which always stores 0.

Thus, we can now have a particular register as stack pointer, so we can have :- (R5 is stack pointer here).

SWI R31, (R5)0

ADD R5, R5, 4



Q. Write code to sum the elements of  $A[100]$ , where start location of the array is  $0x00001000$ .

```

→ ADD R1, R0, 100          ← 0
  ADD R2, R0, 0             ← 4
  ADD R3, R0, 0x00001000    ← 8
  LOOP:                     ← 12
    LWI R4, (R3)0
    ADD R2, R2, R4          ← 16
    ADD R3, R3, 4           ← 20
    SUB R1, R1, 1           ← 24
    BNEQ R1, LOOP           ← 28
    SWI SWI R2, (R3)0      ← 32
    HLT                     ← 36
  
```

So, Loop translates to  $-\frac{(32-12)}{4} = -5$  at instruction at address 28.

\* In a function call, when we pass variables, they are stored in the stack. Then, as needed, they can be stored in registers. This also necessitates preserving the registers on the stack.

\* ARM ISA:

- Predicated execution: Eg:- `ADDZ R1, R1, #1`  
 $\uparrow$  if zero flag is set.




The flags can be set by instructions suffixed by S, like `MOVS`

- Programmable program counter: R15 is PC. You can carry out a jump like `ADD R15, R15, R1`.
- Instructions to push and pop all registers to/from memory.




## \* MACHINE ORGANISATION FOR RISC (MIPS):


Simple instructions, few instruction formats, fixed encoding



Let us assume that all ALU ops are R type, all memory addressing is base+displacement (I type), BEQ is I type and Jump is J type.



CPU looks like:



The data path consists of:-

- 1) Functional units : Eg: ALU
- 2) Storage : Programmer's Registers, Program counter, temporary registers, cache
- 3) Steering logic (communication infrastructure):  
 It can be broadcast based (IO Bus + Tri-state logic), or point to point (Multiplexer based).

The control unit can either be based on combinational logic or finite state machine.


The state of the machine as visible to the programmer is given by memory+registers. State transitions can also occur in one clock cycle or multiple clock cycles.



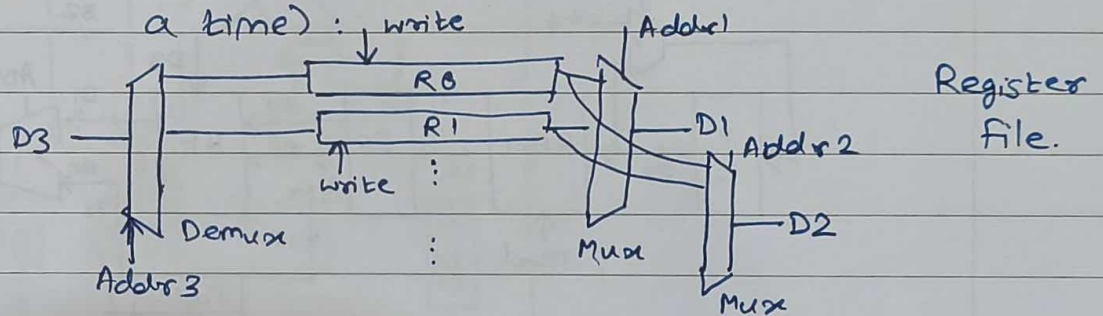
For the multiple cycle model, the control unit has to be designed as a finite state machine (where the internal states aren't visible to the programmer).

### \* DESIGN :-


- Functional units : ALU




- Storage : ~~Memory~~ : Programmer's registers (can read only 2 at a time) :



Program counter : Single register :



- Memory :-



### \* A simple instruction in execution :-

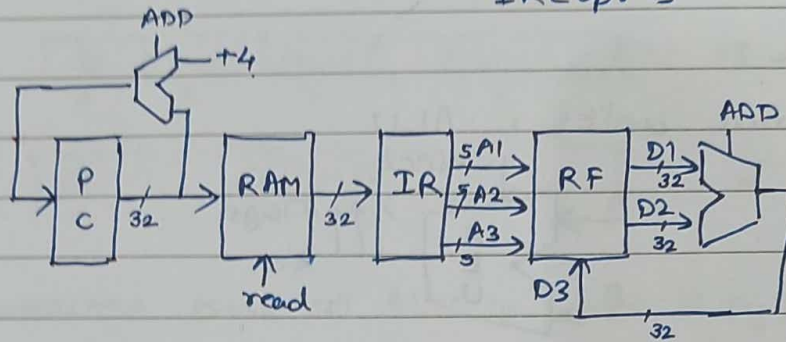
- R Type : ADD R1, R2, R3 :

Micro ~~code~~ code :-

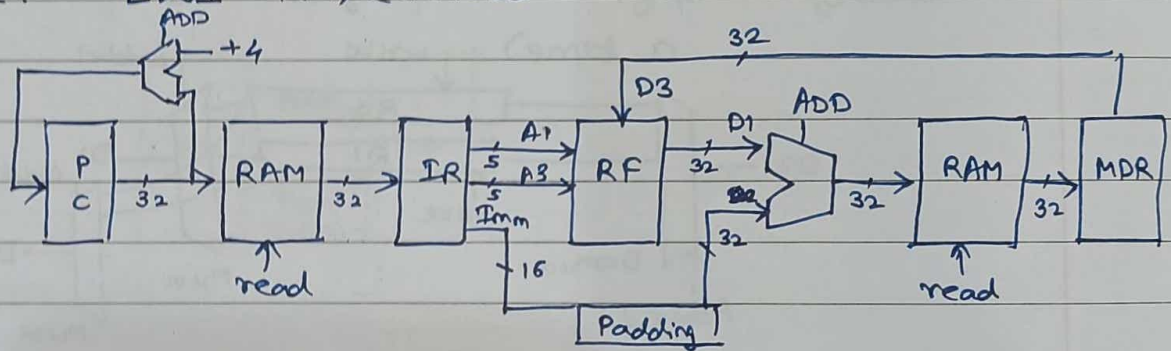
- PC  $\rightarrow$  Memory  $\rightarrow$  IR
- Increment PC by 4 : PC  $\rightarrow$  alu A, 4  $\rightarrow$  alu B, alu C  $\rightarrow$  PC
- IR[operation]  $\rightarrow$  Control unit



- $IR[Op1] \rightarrow \text{Register file } A_1 \rightarrow D1 \rightarrow \text{Alu}_a$
- $IR[Op2] \rightarrow \text{Register file } A_2 \rightarrow D2 \rightarrow \text{Alu}_b$
- $\text{Add in ALU} \rightarrow \text{Alu}_c \rightarrow D3 \text{ of RF} \rightarrow \text{Register file}$   
 $IR[Op3] \rightarrow$



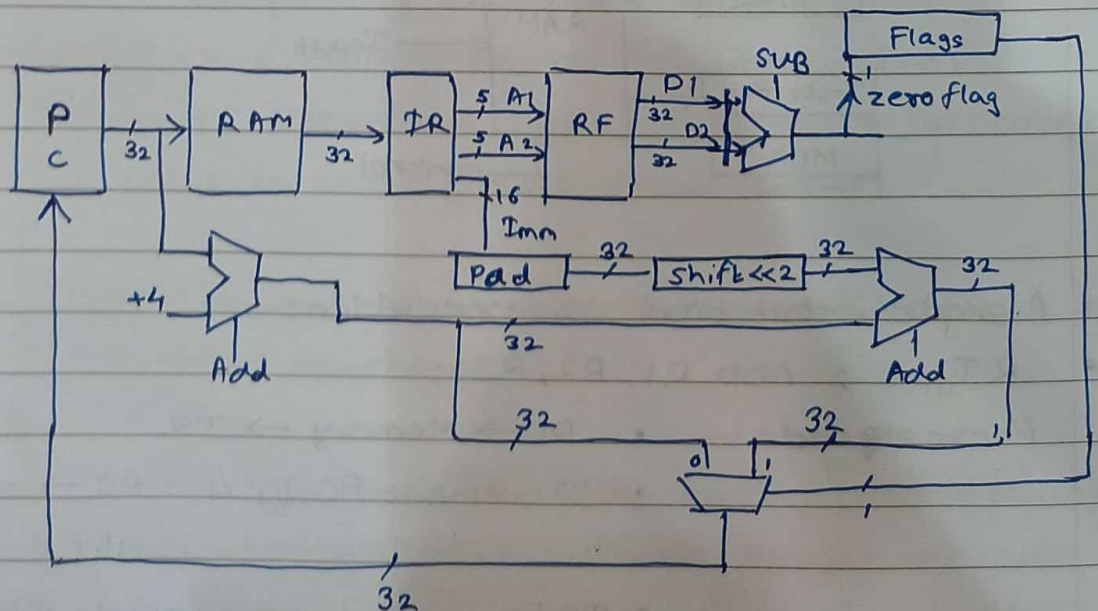
- I Type: `LWI R2, (R1)1000`



`SWI R2, (R1)1000`

Just change read to write and take D2 as input

`BES, R1, R2, 100`





[illegible]

All these devices need control signals, which need to be supplied parallelly/sequentially. The circuit which does this is called the controller. This can be constructed using combinational logic.

The time taken depends on the longest (most delayed) path. The technological bottleneck is not the clock speed, it is the circuit delay.

The total time taken for the program execution

$$= \underbrace{IC}_{\text{Dynamic instruction count}} \times CPI \times \text{Clock time period}$$

Dynamic instruction count: # of instructions actually executed.

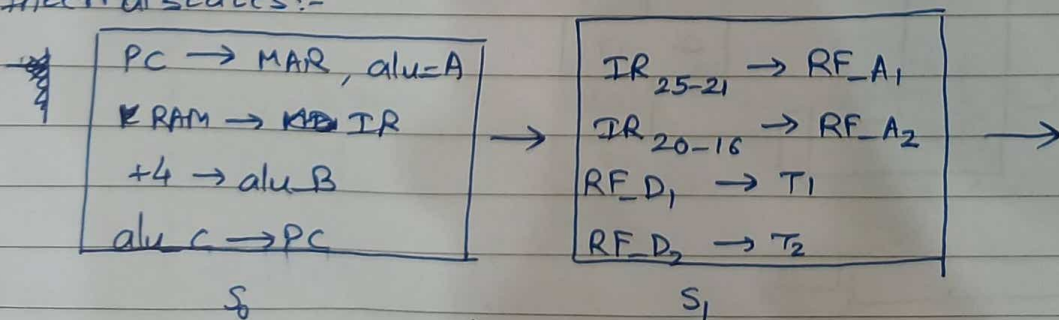
~~Also~~ The CPI can be made 1, but that is expensive in terms of resources. We can instead do it step by step over multiple clock cycles. The subtasks in this process are generally:

- 1) Fetching instruction
- 2) Update PC
- 3) Read operands
- 4) Instruction computation / Memory address
- 5) Read/write memory
- 6) Update the ~~state~~ RF


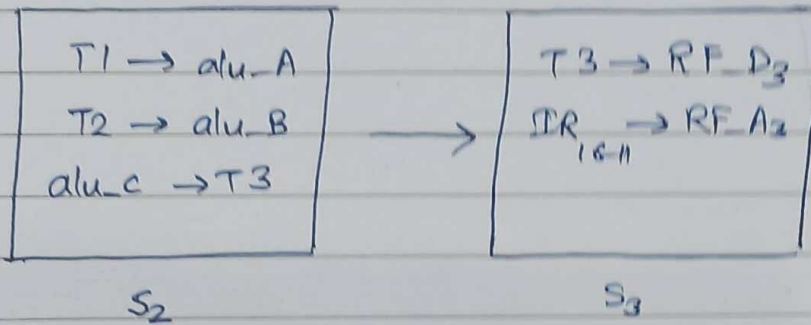
Thus, for an instruction, the system maintains a state for each clock cycle. We also need registers (like the IR, MAR, MDR) at the boundaries between these states.

- R type: 1, 2, 3, 4, 6

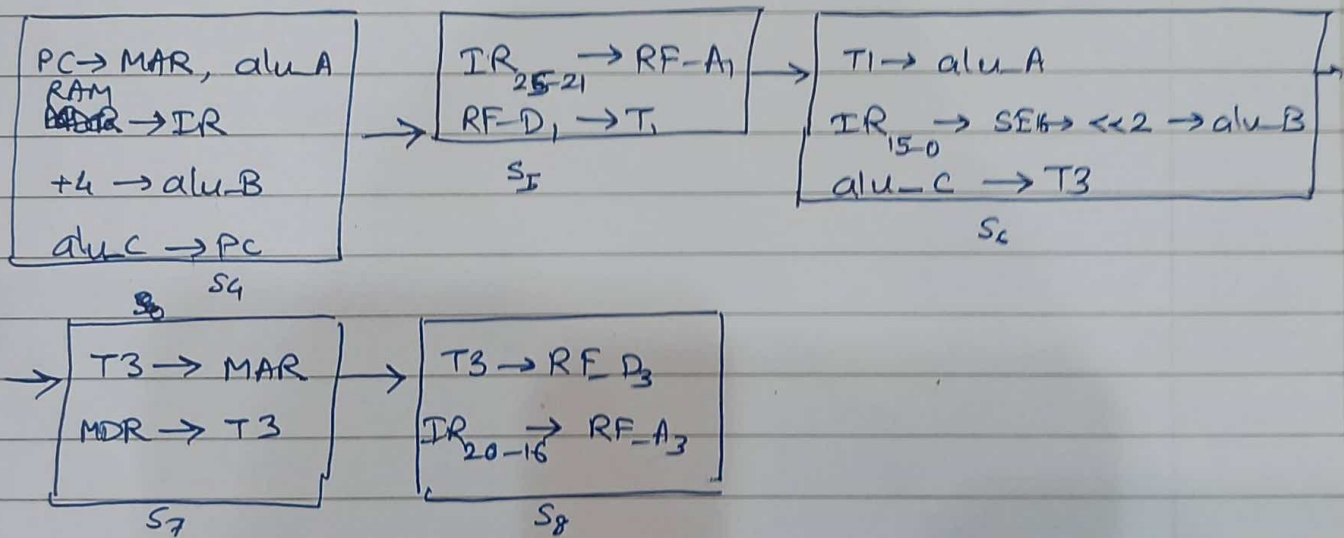
Internal states:-







- I type: 1, 2, 3, 4, 5, 6 (LW)



\* Note: In any state, read happens ~~a little~~ in the beginning. Write happens for all registers at the very end.

- I type: 1, 2, 3, 4, 5 (SW).

