BERT-ASR

Wen-Chin Huang[1,2], Chia-Hua Wu[2], Shang-Bao Luo[2], Kuan-Yu Chen[3], Hsin-Min Wang[2], Tomoki Toda[1]
[1]Nagoya University, Japan [2]Academia Sinica, Taiwan [3]National Taiwan University of Science and Technology, Taiwan

Overview

- A simple method for automatic speech recognition (**ASR**) by fine-tuning BERT
- **BERT-ASR**, formulates ASR as a classification problem, where the objective is to correctly classify the next word given the acoustic speech signals and the history words.

BERT

- **B**idirectional **E**ncoder **R**epresent**a**tions from **T**ransformers
- A language model (LM) trained on large-scale unlabeled text data and can generate rich contextual representations
- Learns language by using 2 main methods: **MLM** (Masked Language Model) & **NSP** (Next Sentence Prediction)
- Adopts a multi-layer Transformer encoder architecture

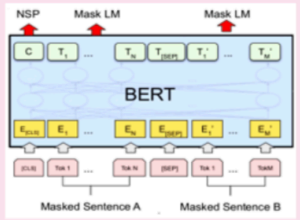

BERT-LM

$$(y_1,\ldots,y_T) \rightarrow \begin{cases} ([\text{CLS}]) \\ ([\text{CLS}], y_1) \\ ([\text{CLS}], y_1, y_2) \\ \ldots \\ ([\text{CLS}], y_1, \ldots, y_{t-1}) \end{cases}$$

- A probabilistic **Language Model (LM)** using BERT
- Exhaustively enumerate all possible training samples
- Training becomes simply minimizing cross-entropy objective

$$\mathcal{L}_{\text{LM}} = -\sum_{i=1}^{N}\sum_{t=1}^{T} P(y_t^{(i)} | [\text{CLS}], y_1^{(i)}, \ldots, y_{t-1}^{(i)}).$$

BERT-ASR Training and Fine-Tuning

- Acoustic frames can be segmented into T groups (T is # of tokens in transcript)
- Acoustic embeddings concatenated with original BERT inputs fed into model.
- This way we augment the BERT-LM into BERT-ASR

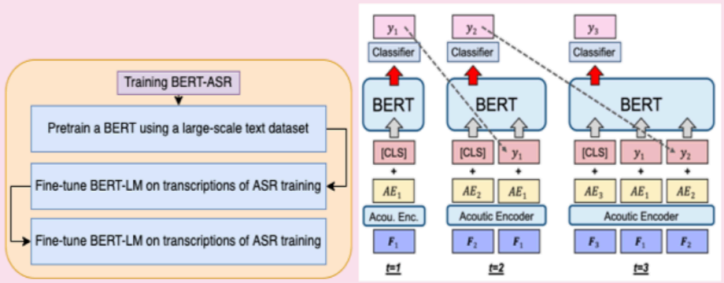

Acoustic Encoder

- Converts raw acoustic feature segments (Fi) into acoustic embeddings
- Authors experimented with the
  - ➢ **Average Encoder** and
  - ➢ **Conv1d-Resnet Encoder** (takes temporal relationship between segments into account)

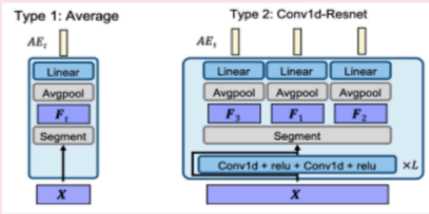

Experiments

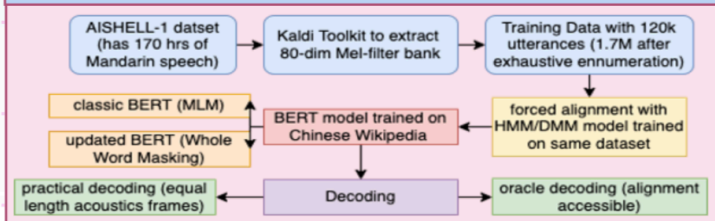

Table 1: Results on the AISHELL-1 dataset. "Orac." and "Prac." denote the oracle decoding and practical decoding, respectively. "Conv1d resnet X" denotes the conv1d resnet encoder with X resnet blocks. Best performance of the BERT-ASR are shown in bold.

Model	Acoustic encoder	Perplexity		CER (Orac.)		CER (Prac.)		SER (Orac.)		SER (Prac.)	
		Dev	Test	Dev	Test	Dev	Test	Dev	Test	Dev	Test
Trigram-LM	-	133.32	127.88	-		-		-		-	
LSTM-LM	-	79.97	78.80	-		-		-		-	
BERT-LM	-	39.74	41.72	-		-		-		-	
BERT-ASR	Average	5.88	9.02	65.8	68.9	96.4	105.8	60.3	63.5	91.5	100.3
	Conv1d resnet 1	4.91	7.63	55.8	59.0	89.6	99.6	50.0	53.8	84.4	94.1
	Conv1d resnet 2	**4.77**	**6.94**	**54.6**	**58.8**	89.7	**99.1**	49.5	53.6	84.6	93.5
	Conv1d resnet 3	4.83	7.41	54.8	58.9	89.8	99.4	49.6	53.6	84.6	93.9
	Conv1d-resnet 4	4.78	7.29	**54.6**	59.0	**89.5**	99.3	49.4	53.9	84.4	93.8
GMM-HMM	-	-		-		10.4	12.2	-		-	
DNN-HMM	-	-		-		7.2	8.4	-		-	

References

[4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," 2019
[5] A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever, "Improving language understanding by generative pre- training," 2018.
[8] A.Vaswani,N.Shazeer,N.Parmar,J.Uszkoreit,L.Jones,A.N Gomez, L. Kaiser, and I. Polosukhin, "Attention is All you Need," *NeruIPS*, 2017.
[17] J. Shin, Y. Lee, and K. Jung, "Effective sentence scoring method using BERT for speech recognition," 2019
[18] J.Salazar,D.Liang,T.Q.Nguyen,andK.Kirchhoff,"Masked language model scoring," 2020
[19] H. Futami, H. Inaguma, S. Ueno, M. Mimura, S. Sakai, and T. Kawahara, "Distilling the knowledge of bert for sequence- to-sequence asr," 2020.
[20] H. Bu, J. Du, X. Na, B. Wu, and H. Zheng, "Aishell-1: An open-source mandarin speech corpus and a speech recognition baseline," 2017
[21] D. Povey, A. Ghoshal, G. Boulianne, L. Burget, O. Glembek, N. Goel, M. Hannemann, P. Motlicek, Y. Qian, P. Schwarz, J. Silovsky, G. Stemmer, and K. Vesely, "The kaldi speech recognition toolkit," Dec. 2011.
[23] T.Wolf,L.Debut,V.Sanh,J.Chaumond,C.Delangue,A.Moi, P. Cistac, T. Rault, R. Louf, M. Funtowicz, et al., "Hugging- face's transformers: State-of-the-art natural language processing," 2019.