


Graphs – Definition

- A **graph** $G = (V, E)$ is composed of:
 - V : set of **vertices**
 - $E \subset V \times V$: set of **edges** connecting the **vertices**
- An **edge** $e = (u, v)$ is a pair of vertices
- (u, v) is ordered, if G is a **directed** graph



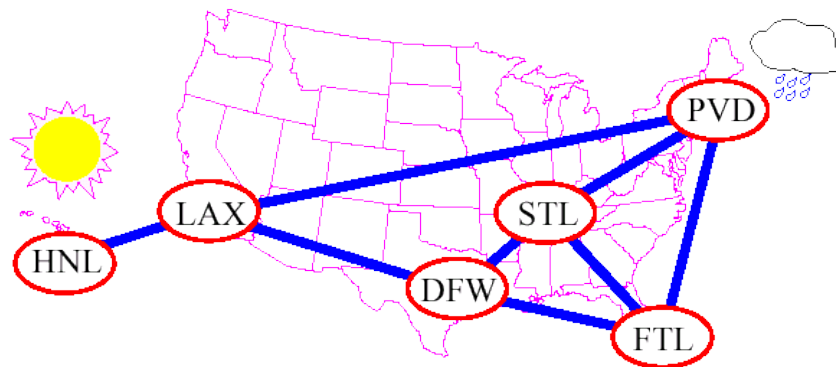
$V = \{a, b, c, d, e\}$

$E =$

$\{(a, b), (a, c), (a, d),$
 $(b, e), (c, d), (c, e),$
 $(d, e)\}$


Applications

- Electronic circuits, pipeline networks
- Transportation and communication networks
- Modeling any sort of relationships (between components, people, processes, concepts)



Graph Terminology

- **adjacent vertices**: connected by an edge
- **degree** (of a **vertex**): # of adjacent vertices




$$\sum_{v \in V} \deg(v) = 2(\# \text{ of edges})$$

Since adjacent vertices each count the adjoining edge, it will be counted twice

- **path**: sequence of vertices v_1, v_2, \dots, v_k such that consecutive vertices v_i and v_{i+1} are adjacent


Graph Terminology (2)

- **simple path:** no repeated vertices




Graph Terminology (3)

- **cycle**: simple path, except that the last vertex is the same as the first vertex




- **connected graph**: any two vertices are connected by some path




Graph Terminology (4)

- **subgraph**: subset of vertices and edges forming a graph
- **connected component**: maximal connected subgraph. E.g., the graph below has 3 connected components




Graph Terminology (5)

- (free) tree - connected graph without cycles
- forest - collection of trees




Data Structures for Graphs

- How can we represent a graph?
 - To start with, we can store the vertices and the edges in two containers, and we store with each edge object references to its start and end vertices





Edge List

- The **edge list**
 - Easy to implement
 - Finding the edges incident on a given vertex is inefficient since it requires examining the entire edge sequence



Adjacency List


- The **Adjacency list** of a vertex v : a sequence of vertices adjacent to v
- Represent the graph by the adjacency lists of all its vertices



$$\text{Space} = \Theta(n + \sum \deg(v)) = \Theta(n + m)$$

Adjacency Matrix

- Matrix M with entries for all pairs of vertices
- $M[i,j] = \text{true}$ – there is an edge (i,j) in the graph
- $M[i,j] = \text{false}$ – there is no edge (i,j) in the graph
- Space = $O(n^2)$



Graph Searching Algorithms

- Systematic search of every edge and vertex of the graph
- Graph $G = (V, E)$ is either directed or undirected
- Today's algorithms assume an adjacency list representation
- Applications
 - Compilers
 - Graphics
 - Maze-solving
 - Mapping
 - Networks: routing, searching, clustering, etc.





Breadth First Search

- A **Breadth-First Search (BFS)** traverses a **connected component** of a graph, and in doing so defines a **spanning tree** with several useful properties
- BFS in an **undirected** graph G is like wandering in a labyrinth with a string.
- The starting vertex s is assigned a distance 0.
- In the first round, the string is unrolled the length of one edge, and all of the edges that are only one edge away from the anchor are visited (**discovered**), and assigned distances of 1





Breadth-First Search (2)

- In the second round, all the new edges that can be reached by unrolling the string 2 edges are visited and assigned a distance of 2
- This continues until every vertex has been assigned a level
- The label of any vertex v corresponds to the length of the shortest path (in terms of edges) from s to v


BFS Example



BFS Example



BFS Example: Result



BFS Algorithm

BFS (G, s)

```
01 for each vertex  $u \in V[G] - \{s\}$ 
02      $\text{color}[u] \leftarrow \text{white}$ 
03      $d[u] \leftarrow \infty$ 
04      $\pi[u] \leftarrow \text{NIL}$ 
```

Init all
vertices

```
05  $\text{color}[s] \leftarrow \text{gray}$ 
06  $d[s] \leftarrow 0$ 
07  $\pi[s] \leftarrow \text{NIL}$ 
08  $Q \leftarrow \{s\}$ 
```

Init BFS with s

```
09 while  $Q \neq \emptyset$  do
10      $u \leftarrow \text{head}[Q]$ 
11     for each  $v \in \text{Adj}[u]$  do
12         if  $\text{color}[v] = \text{white}$  then
13              $\text{color}[v] \leftarrow \text{gray}$ 
14              $d[v] \leftarrow d[u] + 1$ 
15              $\pi[v] \leftarrow u$ 
16              $\text{Enqueue}(Q, v)$ 
17      $\text{Dequeue}(Q)$ 
```

Handle all u 's
children before
handling any
children of
children

```
18      $\text{color}[u] \leftarrow \text{black}$ 
```

BFS Running Time

- Given a graph $G = (V, E)$
 - Vertices are enqueued if their color is white
 - Assuming that en- and dequeuing takes $O(1)$ time the total cost of this operation is $O(V)$
 - Adjacency list of a vertex is scanned when the vertex is dequeued (and only then...)
 - The sum of the lengths of all lists is $\Theta(E)$. Consequently, $O(E)$ time is spent on scanning them
 - Initializing the algorithm takes $O(V)$
- **Total running time $O(V+E)$** (linear in the size of the adjacency list representation of G)

BFS Properties

- Given a graph $G = (V, E)$, BFS **discovers all vertices reachable from a source vertex s**
- It computes the **shortest distance** to all reachable vertices
- It computes a **breadth-first tree** that contains all such reachable vertices
- For any vertex v reachable from s , the path in the breadth first tree from s to v , corresponds to a **shortest path** in G

Breadth First Tree

- Predecessor subgraph of G

$$G_{\pi} = (V_{\pi}, E_{\pi})$$

$$V_{\pi} = \{v \in V : \pi[v] \neq \text{NIL}\} \cup \{s\}$$

$$E_{\pi} = \{(\pi[v], v) \in E : v \in V_{\pi} - \{s\}\}$$

- G_{π} is a breadth-first tree
 - V_{π} consists of the vertices reachable from s , and
 - for all $v \in V_{\pi}$, there is a unique simple path from s to v in G_{π} that is also a shortest path from s to v in G
- The edges in G_{π} are called tree edges

Depth-First Search

- **A depth-first search (DFS)** in an undirected graph G is like wandering in a labyrinth with a **string** and a **can of paint**
 - We start at vertex s , tying the end of our string to the point and painting s “visited (discovered)”. Next we label s as our current vertex called u
 - Now, we travel along an arbitrary edge (u,v) .
 - If edge (u,v) leads us to an already visited vertex v we return to u
 - If vertex v is unvisited, we unroll our string, move to v , paint v “visited”, set v as our current vertex, and repeat the previous steps

Depth-First Search (2)

- Eventually, we will get to a point where **all incident edges on u lead to visited vertices**
- We then **backtrack** by unrolling our string to a previously visited vertex v . Then v becomes our current vertex and we repeat the previous steps
- Then, if all incident edges on v lead to visited vertices, we backtrack as we did before. We **continue to backtrack along the path we have traveled**, finding and exploring unexplored edges, and repeating the procedure

DFS Algorithm

- Initialize – color all vertices white
- Visit each and every white vertex using DFS-Visit
- Each call to DFS-Visit(u) roots a new tree of the depth-first forest at vertex u
- A vertex is **white** if it is undiscovered
- A vertex is **gray** if it has been discovered but not all of its edges have been discovered
- A vertex is **black** after all of its adjacent vertices have been discovered (the adj. list was examined completely)

DFS Algorithm (2)

DFS(G)

```
1 for each vertex  $u \in V[G]$ 
2   do  $color[u] \leftarrow \text{WHITE}$ 
3  $time \leftarrow 0$ 
```

Init all
vertices

```
4 for each vertex  $u \in V[G]$ 
5   do if  $color[u] = \text{WHITE}$ 
6     then DFS-VISIT( $u$ )
```

DFS-VISIT(u)







```
1  $color[u] \leftarrow \text{GRAY}$            ▷ White vertex  $u$  discovered.
2  $d[u] \leftarrow time$              ▷ Mark with discovery time.
3  $time \leftarrow time + 1$          ▷ Tick global time.
```

```
4 for each  $v \in Adj[u]$              ▷ Explore all edges  $(u, v)$ .
5   do if  $color[v] = \text{WHITE}$ 
6     then DFS-VISIT( $v$ )
```







Visit all children
recursively

```
7  $color[u] \leftarrow \text{BLACK}$        ▷ Blacken  $u$ ; it is finished.
8  $f[u] \leftarrow time$              ▷ Mark with finishing time.
9  $time \leftarrow time + 1$          ▷ Tick global time.
```





DFS Example




DFS Example (2)



DFS Example (3)



Assume $G = (V, E)$ is connected
undirected DFS




$G = (V, E)$

T : tree edges


(V, T)

$|T| = n - 1$

Is (V, T) connected?



$T \subseteq E$



all edges of G which are not tree edges
have an ancestor-descendant relationship between their
end-points





$$a \Leftrightarrow b$$

Equivalent definitions of a tree

- a) Connected graph containing no cycles
- b) connected graph containing $|V| - 1$ edges
- c) minimal connected graph

$C \Rightarrow a$







Induction on n of vertices
 every connected graph on k vertices which does not
 have a cycle has $k-1$ edges

Ind step: Consider a graph on $k+1$ vertices which does not have
 a cycle. Such a graph has at least one vertex of
 degree 1 (say w)

~~total of~~

Complete graph has
 k edges w






B: back edges from node to ancestor

F: forward edges: from node to descendant

C: cross edges: from node to another node which is not an ancestor/descendant



DFS Algorithm (3)

- When DFS returns, every vertex u is assigned
 - a discovery time $d[u]$, and a finishing time $f[u]$
- Running time
 - the loops in DFS take time $\Theta(V)$ each, excluding the time to execute DFS-Visit
 - DFS-Visit is called once for every vertex
 - its only invoked on white vertices, and
 - paints the vertex gray immediately
 - for each DFS-visit a loop iterates over all $Adj[v]$
 - the total cost for DFS-Visit is $\Theta(E)$

– **the running time of DFS is $\Theta(V+E)$**

$$\sum_{v \in V} |Adj[v]| = \Theta(E)$$

Predecessor Subgraph

- Define slightly different from BFS


$$G_{\pi} = (V, E_{\pi})$$

$$E_{\pi} = \{(\pi[v], v) \in E : v \in V \text{ and } \pi[v] \neq \text{NIL}\}$$

- The PD subgraph of a depth-first search forms a **depth-first forest** composed of several depth-first trees
- The edges in G_{π} are called tree edges

DFS Timestamping

- The DFS algorithm maintains a monotonically increasing global clock
 - discovery time $d[u]$ and finishing time $f[u]$
- For every vertex u , the inequality $d[u] < f[u]$ must hold




DFS Timestamping

- Vertex u is
 - white before time $d[u]$
 - gray between time $d[u]$ and time $f[u]$, and
 - black thereafter
- Notice the structure throughout the algorithm.
 - gray vertices form a linear chain
 - corresponds to a stack of vertices that have not been exhaustively explored (DFS-Visit started but not yet finished)

DFS Parenthesis Theorem


- Discovery and finish times have parenthesis structure
 - represent discovery of u with left parenthesis " $(u$ "
 - represent finishing of u with right parenthesis " $u)$ "
 - history of discoveries and finishings makes a well-formed expression (parenthesis are properly nested)
- Intuition for proof: any two intervals are either disjoint or enclosed
 - Overlapping intervals would mean finishing ancestor, before finishing descendant or starting descendant without starting ancestor

DFS Parenthesis Theorem (2)




DFS Edge Classification

- Tree edge (gray to white)
 - encounter new vertices (white)
- Back edge (gray to gray)
 - from descendant to ancestor




DFS Edge Classification (2)

- Forward edge (gray to black)
 - from ancestor to descendant
- Cross edge (gray to black)
 - remainder – between trees or subtrees




DFS Edge Classification (3)

- Tree and back edges are important
- Most algorithms do not distinguish between forward and cross edges



Directed Acyclic Graphs

- A DAG is a directed graph with no cycles



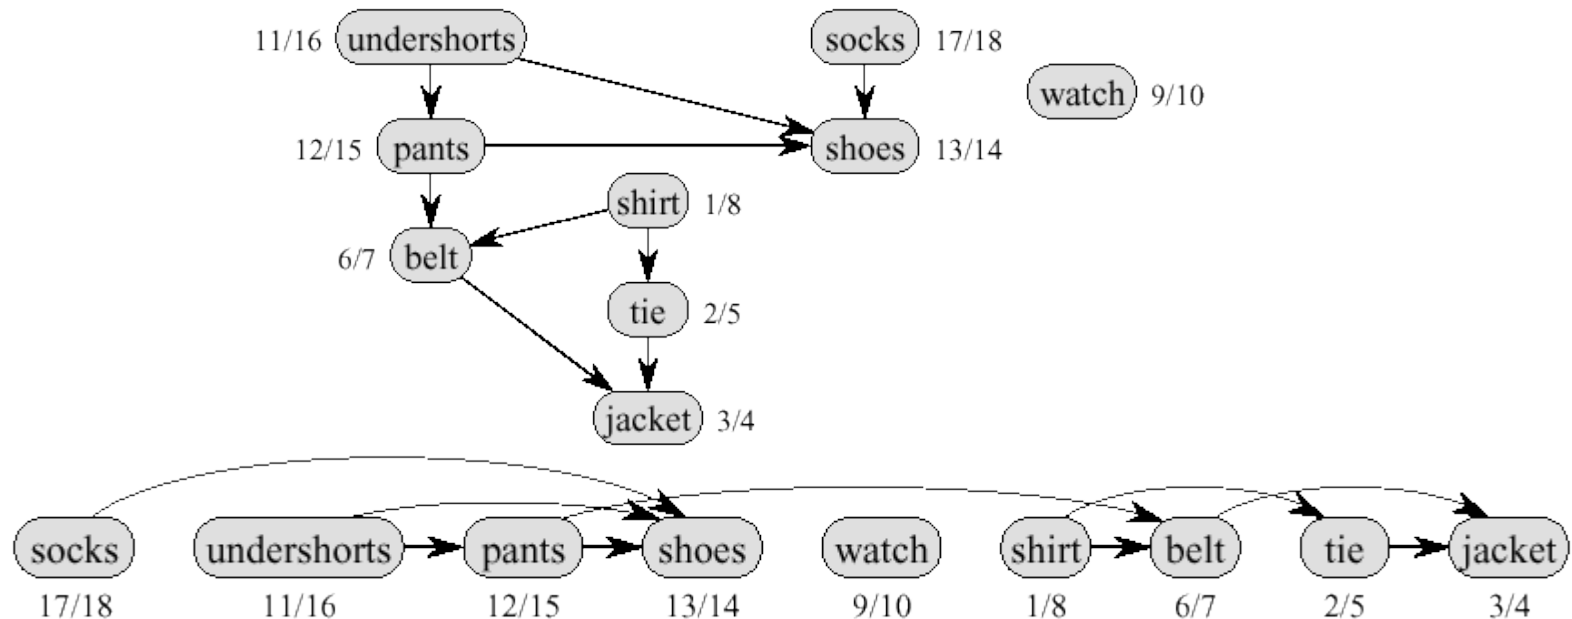
- Often used to indicate precedences among events, i.e., event a must happen before b
- An example would be a parallel code execution
- Total order can be introduced using **Topological Sorting**

DAG Theorem

- A directed graph G is acyclic if and only if a DFS of G yields no back edges. Proof:
 - **suppose there is a back edge (u,v)** ; v is an ancestor of u in DFS forest. Thus, there is a path from v to u in G and (u,v) completes the cycle
 - **suppose there is a cycle c** ; let v be the first vertex in c to be discovered and u is a predecessor of v in c .
 - Upon discovering v the whole cycle from v to u is white
 - We must visit all nodes reachable on this white path before return DFS-Visit(v), i.e., vertex u becomes a descendant of v
 - Thus, (u,v) is a back edge
- Thus, we can verify a DAG using DFS!

Topological Sort Example

- Precedence relations: an edge from x to y means one must be done with x before one can do y
- Intuition: can schedule task only when all of its subtasks have been scheduled



Topological Sort

- Sorting of a directed acyclic graph (DAG)
- A topological sort of a DAG is a linear ordering of all its vertices such that for any edge (u,v) in the DAG, u appears before v in the ordering
- The following algorithm topologically sorts a DAG

Topological-Sort(G)

- 1) call DFS(G) to compute finishing times $f[v]$ for each vertex v
 - 2) as each vertex is finished, insert it onto the front of a linked list
 - 3) return the linked list of vertices
- The linked list comprises a total ordering

Topological Sort

- Running time
 - depth-first search: $O(V+E)$ time
 - insert each of the $|V|$ vertices to the front of the linked list: $O(1)$ per insertion
- Thus the total running time is $O(V+E)$

Topological Sort Correctness

- Claim: for a DAG, an edge $(u,v) \in E \Rightarrow f[u] > f[v]$
- When (u,v) explored, u is gray. We can distinguish three cases
 - $v = \text{gray}$
 $\Rightarrow (u,v) = \text{back edge (cycle, contradiction)}$
 - $v = \text{white}$
 $\Rightarrow v$ becomes descendant of u
 $\Rightarrow v$ will be finished before u
 $\Rightarrow f[v] < f[u]$
 - $v = \text{black}$
 $\Rightarrow v$ is already finished
 $\Rightarrow f[v] < f[u]$
- The definition of topological sort is satisfied